

Experimental Pipeline: 41,832 Total Inference Calls

498 Questions
(300 MMLU + 198 GPQA)

**12 Open-Weight
Reasoning Models**

Baseline Runs (no hint)
5,976 inference calls



6 Hint Types Injected
(sycophancy, consistency, visual, metadata, grader, unethical)



Hinted Runs
35,856 inference calls



Filter: answer changed to match hint?
10,276 influenced cases (28.7%)

Regex + 3-Judge
LLM Pipeline
(GLM-5, Kimi K2, Gemini 3)



82.6% faithful

Claude Sonnet 4
(independent judge)
\$48.99 total



69.7% faithful